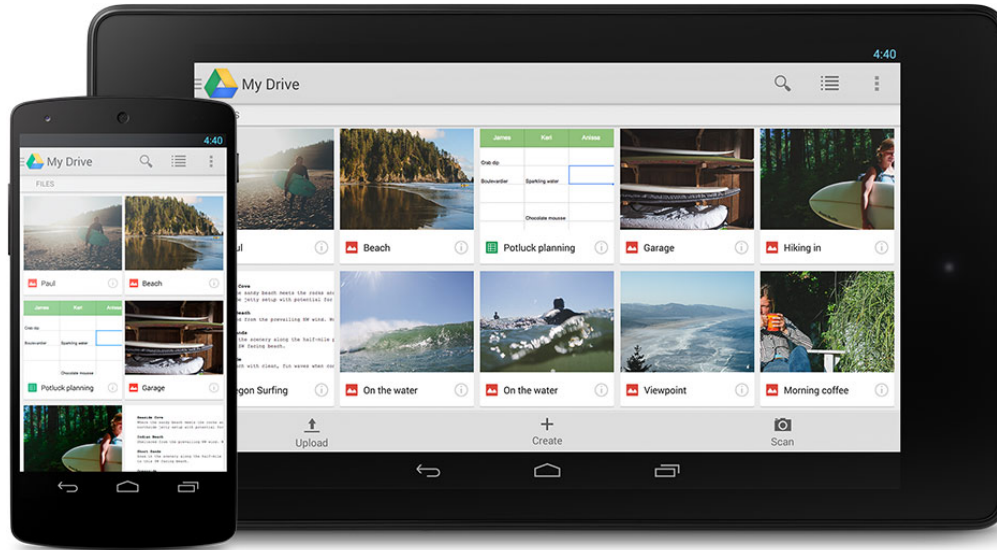# Large-scale Data Systems

Lecture 4: Shared memory

Prof. Gilles Louppe
g.louppe@uliege.be

LIÈGE université

Today

Towards a distributed file system.

- How do you share resources between processes?

- Can we build the illusion of single storage?

  - While replicating data for fault-tolerance and scalability?

  - While maintaining consistency?

Real shared memory


```
Processor 1    Processor 2   ....   Processor n
```

Shared main memory modules / I/O units
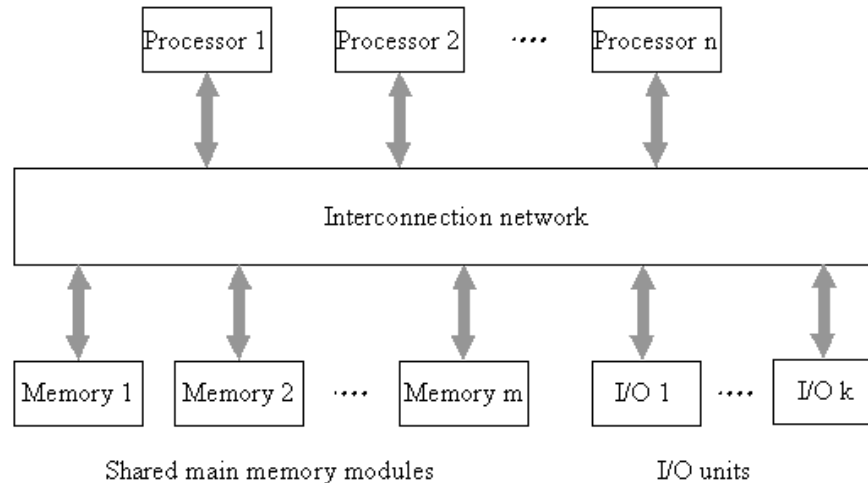
In a multiprocessor machine, processors typically communicate through shared memory provided at the hardware level (e.g., shared blocks of RAM that can be accessed by distinct CPUs).

- Shared memory can be viewed as an array of registers to which processors can read or write.

- Shared memory systems are easy to program since all processors share a single view of the data.

Shared memory emulation

We want to simulate a shared memory abstraction in a distributed system, on top of message passing communication.

Why?

- Enable shared memory algorithms without being aware that processes are actually communicating by exchanging messages.
    - This is often much easier to program.
- Equivalent to consistent data replication across nodes.

Data replication

- Why replicating data across nodes? Shared data allows to:

  - Reduce network traffic

  - Promote increased parallelism

  - Be robust against failures

  - Result in fewer page faults

- Applications:

  - distributed databases

  - distributed file systems

  - distributed cache

  - ...

- Challenges:

  - Consistency in presence of failures.

  - Consistency in presence of concurrency.

Regular registers

Read/Write registers

- A register represents each memory location.

- A register contains only positive integers and is initialized to 0.

- Registers have two operations:
  - $read()$: return the current value of the register.
  - $write(v)$: update the register to value v.

- An operation is not instantaneous:
  - It is first invoked by the calling process.
  - It computes for some time.
  - It returns a response upon completion.

Definitions

- In an execution, an operation is

    - completed if both invocation and response occurred.

    - failed if invoked but not no response was received.

- Operation o_1 precedes o_2 if response of o_1 precedes the invocation of o_2.

- Operations o_1 and o_2 are concurrent if neither precedes the other.

- $(1, N)$ register: 1 designated writer, N readers.

- (M, N) register: M writers, N readers.

Regular registers ($onrr$)

Module:

    Name: $(1, N)$-RegularRegister, **instance** *onrr*.

Events:

    Request: ⟨ *onrr*, *Read* ⟩: Invokes a read operation on the register.

    Request: ⟨ *onrr*, *Write* | *v* ⟩: Invokes a write operation with value *v* on the register.

    Indication: ⟨ *onrr*, *ReadReturn* | *v* ⟩: Completes a read operation on the register with return value *v*.

    Indication: ⟨ *onrr*, *WriteReturn* ⟩: Completes a write operation on the register.

Properties:

    ONRR1: *Termination:* If a correct process invokes an operation, then the operation eventually completes.

    ONRR2: *Validity:* A read that is not concurrent with a write returns the last value written; a read that is concurrent with a write returns the last value written or the value concurrently written.

Regular register example (1)


p1 ---- w(5) ---- w(6) ---->

p2 ---- r() → 0 ---- r() → 5 ---- r() → 6 ---->

[Q] Regular or non-regular?

Regular register example (2)

[Q] Regular or non-regular?

Regular register example (3)

[Q] Regular or non-regular?

Centralized algorithm

- Designates one process as the leader.

    - E.g., using the leader election abstraction (see Lecture 2).

- To $read()$:

    - Ask the leader for latest value.

- To $write(v)$:

    - Update leader's value to v.

[Q] Problem? Does not work if leader crashes!

Decentralized algorithm (bogus)

- Intuitively, make an algorithm in which

  - A $\mathrm{read}()$ reads the local value.

  - A $\mathrm{write}(v)$ writes to all nodes.

- To $\mathrm{read}()$:

  - Return local value.

- To $\mathrm{write}(v)$:

  - Update local value to v.

  - Broadcast v to all (each node then locally updates).

  - Return.

[Q] Problem?

Decentralized algorithm (bogus) example

p1 ------●──w(5)──●--➤

p2 ----------------●──r() → 0──●─●──r() → 0──●--➤●──r() → 5──●----➤

Validity is violated!

Read-one Write-all algorithm

- Bogus algorithm modified.

- To $\mathrm{read}()$:

    - Return local value.

- To $\mathrm{write}(v)$:

    - Update local value to v.

    - Broadcast v to all (each node locally updates).

    - Wait for acknowledgement from all correct nodes.

        - Require a perfect failure detector (fail-stop).

    - Return.

Implements:
    $(1, N)$-RegularRegister, **instance** *onrr*.

Uses:
    BestEffortBroadcast, **instance** *beb*;
    PerfectPointToPointLinks, **instance** *pl*;
    PerfectFailureDetector, **instance** $\mathcal{P}$.

upon event $\langle$ *onrr, Init* $\rangle$ **do**
    $val := \bot$;
    $correct := \Pi$;
    $writeset := \emptyset$;

upon event $\langle$ $\mathcal{P}$, *Crash* $\mid p$ $\rangle$ **do**
    $correct := correct \setminus \{p\}$;

upon event $\langle$ *onrr, Read* $\rangle$ **do**
    trigger $\langle$ *onrr, ReadReturn* $\mid val$ $\rangle$;

upon event $\langle$ *onrr, Write* $\mid v$ $\rangle$ **do**
    trigger $\langle$ *beb, Broadcast* $\mid [\text{WRITE}, v]$ $\rangle$;

upon event $\langle$ *beb, Deliver* $\mid q, [\text{WRITE}, v]$ $\rangle$ **do**
    $val := v$;
    trigger $\langle$ *pl, Send* $\mid q, \text{ACK}$ $\rangle$;

**upon event** $\langle$ *pl, Deliver* $\mid p, \text{ACK}$ $\rangle$ **do**
    $writeset := writeset \cup \{p\}$;

upon $correct \subseteq writeset$ **do**
    $writeset := \emptyset$;
    trigger $\langle$ *onrr, WriteReturn* $\rangle$;

Read-one Write-all example

Validity is no longer violated because the write response has been postponed.

# Quorum principle

Can we implement a regular register in fail-silent? (without a failure detector)

Quorum principle

Can we implement a regular register in fail-silent? (without a failure detector)

- Assume a majority of correct nodes.

- Divide the system into two overlapping majority quorums.
  - i.e., each quorum counts at least $\lfloor \frac{N}{2} \rfloor + 1$ nodes.

- Always write to and read from a majority of nodes.

- At least one node must know the most recent value.

Majority voting algorithm

Implements:
    $(1, N)$-RegularRegister, **instance** *onrr*.
Uses:
    BestEffortBroadcast, **instance** *beb*;
    PerfectPointToPointLinks, **instance** *pl*.

upon event $\langle$ *onrr, Init* $\rangle$ **do**
    $(ts, val) := (0, \bot)$;
    $wts := 0$;
    $acks := 0$;
    $rid := 0$;
    $readlist := [\bot]^N$;

upon event $\langle$ *onrr, Write* $\mid v$ $\rangle$ **do**
    $wts := wts + 1$;
    $acks := 0$;
    trigger $\langle$ *beb, Broadcast* $\mid$ [WRITE$, wts, v$] $\rangle$;

upon event $\langle$ *beb, Deliver* $\mid p,$ [WRITE$, ts', v'$] $\rangle$ **do**
    if $ts' > ts$ **then**
        $(ts, val) := (ts', v')$;
    trigger $\langle$ *pl, Send* $\mid p,$ [ACK$, ts'$] $\rangle$;

upon event $\langle$ *pl, Deliver* $\mid q,$ [ACK$, ts'$] $\rangle$ **such that** $ts' = wts$ **do**
    $acks := acks + 1$;
    if $acks > N/2$ **then**
        $acks := 0$;
        trigger $\langle$ *onrr, WriteReturn* $\rangle$;

upon event $\langle$ *onrr, Read* $\rangle$ **do**
    $rid := rid + 1$;
    $readlist := [\bot]^N$;
    trigger $\langle$ *beb, Broadcast* $\mid$ [READ, *rid*] $\rangle$;

upon event $\langle$ *beb, Deliver* $\mid p$, [READ, *r*] $\rangle$ **do**
    trigger $\langle$ *pl, Send* $\mid p$, [VALUE, *r, ts, val*] $\rangle$;

upon event $\langle$ *pl, Deliver* $\mid q$, [VALUE, r, ts', v'] $\rangle$ **such that** $r = rid$ **do**
    $readlist[q] := (ts', v')$;
    if $\#(readlist) > N/2$ **then**
        $v := \text{highestval}(readlist)$;
        $readlist := [\bot]^N$;
        trigger $\langle$ *onrr, ReadReturn* $\mid v$ $\rangle$;

[Q] Why resetting `acks` and `readlist` right after having received back just more than $N/2$ messages?

Atomic registers

Towards single storage illusion.

Sequential consistency

An operation o_1 locally precedes o_2 in E if o_1 and o_2 occur at the same node and o_1 precedes o_2 in E.

An execution E is sequentially consistent if an execution F exists such that:

- E and F contain the same events;

- F is sequential;

- Read responses have value of the preceding write invocation in F;

- If o_1 locally precedes o_2 in E, then o_1 locally precedes o_2 in F.

Example (1)

Sequential consistency disallows such execution.

Example (2)

p1   w(5)

p2   r() → 0

p3   r() → 5

happened-before
observable

Sequential consistency allows such execution.

Linearization

- **Linearizability**:
  - Read operations appear as if immediately happened at all nodes at time between invocation and response.
  - Write operations appear is if immediately happened at all anode at time between invocation and response.
  - Failed operations appear as
    - completed at every node, XOR
    - never occurred at any node.
  - The hypothetical serial execution is called a linearization of the actual execution.

- **Termination**:
  - If node is correct, each read and write operation eventually completes.

Example (1)

Linearizability disallows such execution.

Example (2)

Linearizability disallows such execution.

$(1, N)$ atomic registers

Module:

   Name: $(1, N)$-AtomicRegister, **instance** *onar*.

Events:

   Request: $\langle$ *onar*, *Read* $\rangle$: Invokes a read operation on the register.

   Request: $\langle$ *onar*, *Write* $| v \rangle$: Invokes a write operation with value v on the register.

   Indication: $\langle$ *onar*, *ReadReturn* $| v \rangle$: Completes a read operation on the register with return value v.

   Indication: $\langle$ *onar*, *WriteReturn* $\rangle$: Completes a write operation on the register.

Properties:

   ONAR1–ONAR2: Same as properties ONRR1–ONRR2 of a $(1, N)$ regular register (Module 4.1).

   ONAR3: *Ordering:* If a read returns a value v and a subsequent read returns a value w, then the write of w does not precede the write of v.

[Q] Show that linearizability is equivalent to validity + ordering.

Atomic register example (1)

[Q] Atomic? No, not possible to find linearization points.

Atomic register example (2)

[Q] Atomic? Yes

Atomic register example (3)

[Q] Atomic? Yes

Regular but not atomic

[Q] Atomic? No. Regular? Yes, using majority voting.

Implementation of $(1, N)$ atomic registers

- When reading, write back the value that is about to be returned.

- Maintain a local timestamp ts and its associated value val.

- Overwrite the local pair only upon a write operation of a more recent value.

Read-Impose Write-all algorithm

Implements:
    $(1, N)$-AtomicRegister, **instance** *onar*.

Uses:
    BestEffortBroadcast, **instance** *beb*;
    PerfectPointToPointLinks, **instance** *pl*;
    PerfectFailureDetector, **instance** $\mathcal{P}$.

upon event $\langle$ *onar, Init* $\rangle$ **do**
    $(ts, val) := (0, \bot)$;
    $correct := \Pi$;
    $writeset := \emptyset$;
    $readval := \bot$;
    $reading := \text{FALSE}$;

upon event $\langle$ $\mathcal{P}$, *Crash* $\mid p$ $\rangle$ **do**
    $correct := correct \setminus \{p\}$;

upon event $\langle$ *onar, Read* $\rangle$ **do**
    $reading := \text{TRUE}$;
    $readval := val$;
    trigger $\langle$ *beb, Broadcast* $\mid [\text{WRITE}, ts, val]$ $\rangle$;

upon event ⟨ *onar, Write* | *v* ⟩ **do**
    trigger ⟨ *beb, Broadcast* | [WRITE, *ts* + 1, *v*] ⟩;

upon event ⟨ *beb, Deliver* | *p*, [WRITE, *ts'*, *v'*] ⟩ **do**
    if *ts'* > *ts* **then**
            (*ts, val*) := (*ts', v'*);
    trigger ⟨ *pl, Send* | *p*, [ACK] ⟩;

upon event ⟨ *pl, Deliver* | *p*, [ACK] ⟩ **then**
    writeset := *writeset* ∪ {*p*};

upon *correct* ⊆ *writeset* **do**
    writeset := ∅;
    if *reading* = TRUE **then**
            reading := FALSE;
            trigger ⟨ *onar, ReadReturn* | *readval* ⟩;
    else
            trigger ⟨ *onar, WriteReturn* ⟩;

[Q] How to adapt to fail-silent? Read-Impose Write-Majority

Correctness

- Ordering: if a read returns v and a subsequent read returns w, then the write of w does not precede the write of v.

  - p writes v with timestamp ts_v.

  - p writes w with timestamp $ts_w > ts_v$.

  - q reads the values the w.

  - some time later, r invokes a read operation.

  - when q completes its read, all correct processes have a timestamp $ts \geq ts_w$.

  - there is no way for r to changes its value back to v after this because $ts_v < ts_w$.

[Q] Show that the termination and validity properties are satisfied.

(N, N) atomic registers

Module:

    Name: (N, N)-AtomicRegister, **instance** *nnar*.

Events:

    Request: $\langle$ *nnar, Read* $\rangle$: Invokes a read operation on the register.

    Request: $\langle$ *nnar, Write* $\mid v$ $\rangle$: Invokes a write operation with value v on the register.

    Indication: $\langle$ *nnar, ReadReturn* $\mid v$ $\rangle$: Completes a read operation on the register with return value v.

    Indication: $\langle$ *nnar, WriteReturn* $\rangle$: Completes a write operation on the register.

Properties:

    NNAR1: *Termination:* Same as property ONAR1 of a $(1, N)$ atomic register (Module 4.2).

    NNAR2: *Atomicity:* Every read operation returns the value that was written most recently in a hypothetical execution, where every failed operation appears to be complete or does not appear to have been invoked at all, and every complete operation appears to have been executed at some instant between its invocation and its completion.

- How do we handle multiple writers?

- Read-Impose Write-all does not support multiple writers:

    - Assume p and q both store the same timestamp ts (e.g., because of a preceding completed operation).

    - When p and q proceed to write, different values would become associated with the same timestamp.

- Fix:

    - Together with the timestamp, pass and store the identity pid of the process that writes a value v.

    - Determine which message is the latest

        - by comparing timestamps,

        - by breaking ties using the process IDs.

[Q] How many messages are exchanged per read and write operations?

[Q] Can we similarly fix Read-Impose Write-Majority?

Simulating message passing?

- As we saw, we can simulate shared with message passing.

  - A majority of correct nodes is all that is needed.

- Can we simulate message passing in shared memory?

  - Yes: use one register pq for every channel.

    - Modeling a directed channel from p to q.

  - Send messages by appending to the right channel.

  - Receive messages by busy-polling incoming "channels".

- Shared memory and message passing are equivalent.

Summary

- Shared memory registers form a shared memory abstraction with read and write operations.

  - Consistency of the data is guaranteed, even in the presence of failures and concurrency.

- Regular registers:

  - Bogus algorithm (does not work)

  - Centralized algorithm (if no failures)

  - Read-One Write-All algorithm (fail-stop)

  - Majority voting (fail-silent)

- Atomic registers:

  - Single writers

  - Multiple writers

The end.